

SALESFORCE AGENTFORCE — INSTRUCTOR REFERENCE GUIDE

Agentforce Observability: From Setup to Analytics

Covers: Setup Toggles · Feature Comparison · Overview · Quality · Optimization · Service Agent Tabs · Tagging

CHAPTER 1

Setup — Observability Toggles

Setup → Einstein & AI → Agentforce Observability — master controls for every layer of visibility

What this screen is: The master control panel for observability across all Agentforce agents and Einstein AI features. Each toggle activates a distinct layer of visibility — from raw audit logs to real-time health alerts. Layers are cumulative: enabling a higher-level toggle often auto-enables the ones below it.

OFF

Audit and Feedback

"Collect and store Einstein generative AI audit and feedback data for analysis and reporting."

WHAT IT IS Activates collection and storage of every Einstein AI interaction — what was asked, what was returned, and any feedback users gave on response quality.

PURPOSE Creates a permanent, queryable paper trail of AI activity in the org — the foundation for compliance and quality investigation.

WHY IT MATTERS Without this on, you have zero record of what the AI did. With it on, you can investigate complaints, prove compliance, and identify patterns in poor outputs.

REAL-WORLD USE CASE

A customer escalates: "the AI gave me wrong information last Tuesday." With Audit on, you pull that exact session and see precisely what prompt was sent and what the model returned.

OFF — NESTED

Knowledge / RAG Quality Data and Metrics

"Monitor retrievers to track run-time performance, view trends, and identify opportunities to improve results."

WHAT IT IS Monitoring layer specifically for the **retrieval-augmented generation (RAG)** pipeline — the step where the agent searches Knowledge articles before generating a response.

WHY IT MATTERS RAG quality is the #1 variable in agent answer quality. Most "hallucinations" are actually retrieval failures — the model answered confidently using the wrong article.

REAL-WORLD USE CASE

Agent answers keep being slightly off. RAG Metrics reveal the retriever ranks a stale 2021 Knowledge article above the updated 2025 version. Fix the article metadata — fix the answers.

RETIRED — SUMMER '26

Prompt Builder Usage and Feedback Metrics

"Prompt Performance Metrics retired in Summer '26 and can't be re-enabled."

WHAT IT WAS Tracked performance of prompts built in Prompt Builder — which prompts performed well, which generated poor outputs, and usage volume over time.

Instructor note — tombstone: This feature is gone as of Summer '26 and cannot be re-enabled. Point it out explicitly — students will encounter it in older orgs and outdated documentation. Its presence is a reminder that the platform evolves fast and feature availability is release-dependent.

ON

Agentforce Session Tracing

"Capture and store detailed interaction data for all agents — reasoning engine executions, actions, prompt/gateway inputs and outputs, error messages, and agent responses."

WHAT IT IS The **master toggle** for Agentforce observability. Captures the complete lifecycle of every agent session — reasoning steps, actions taken, prompt inputs/outputs, error messages, and final responses.

PURPOSE Gives you a full "replay" of what any agent did from start to finish, including every decision the reasoning engine made. Provisions the Session Tracing Data Model (STDM) in Data Cloud — the foundation all other analytics depend on.

WHY IT MATTERS Without this, agents are a black box. Required to unlock Agent Analytics, Agent Optimization, and Agent Health Monitoring. Enabling it automatically turns on Agent Optimization.

REAL-WORLD USE CASE

An agent escalated a case it should have resolved. Session Tracing lets you replay the exact reasoning step that triggered the escalation — was it a missing topic, a misconfigured action, or a bad retrieval result?

OFF — DEVELOPER TOOL

Agent Platform Tracing

"Deep insights into every Platform Action — Flow, Apex — within an agent's workflow. Data stored in Data Cloud."

WHAT IT IS Captures OpenTelemetry-based spans for every instrumented service in an agent transaction: planner steps, LLM calls, knowledge references, Prompt Builder templates, Flow executions, and Apex invocations. ~2-minute SLA to Data Cloud.

KEY DISTINCTION Session Tracing = what the agent *decided*. Platform Tracing = what the platform *did* in response — exact durations, execution paths, errors. This is the engineer's tool for finding bottlenecks.

REQUIREMENTS Data Cloud provisioned + PlatformObservability org permission + Access Agentforce Optimization permission set. GA as of April 2026.

REAL-WORLD USE CASE

Agent responses take 8 seconds. Platform Tracing shows the Apex action it called took 7.2 of those seconds. The AI reasoning was fast — the bottleneck was entirely in custom code.

ON — AUTO-ENABLED

Agent Optimization

"Identify gaps in agent performance — categorizes user interactions by intent to spot trends and understand user needs at scale."

WHAT IT IS AI-powered analysis layer that categorizes interactions by **user intent** at scale. Powers the Insights, Sessions & Intents, and Health Monitoring views in Agentforce Studio.

NEW METRICS Spring '26 additions: User Feedback (Positive/Negative), Task Resolution Rate, Instruction Adherence Rate, Average Toxicity Score, Session Outcome (abandonment/deflection), and Custom Scorers (Beta).

AUTO-ON Turns on automatically when Session Tracing is enabled — it cannot function without the session data Session Tracing captures.

REAL-WORLD USE CASE

5,000 sessions this week. Optimization flags 340 show users re-asking the same question in different ways — a clear signal the agent isn't resolving the intent on first try. Now you know exactly which topic to fix.

OFF — NESTED

Agent Health Monitoring

"Track agent performance in near-real time with metrics, reporting, alerts, and trace logs."

WHAT IT IS Near-real-time performance tracking with configurable metrics, dashboards, alerts, and trace logs for all agents in the org.

KEY DISTINCTION The other features above are **forensic** — they tell you what happened. Health Monitoring is **operational** — it tells you what's happening right now and alerts before users notice.

REAL-WORLD USE CASE

A bad deployment at 2am spikes agent error rates. Health Monitoring fires an alert before a single user calls in. The on-call team rolls back and most customers never notice.

LAYERED OBSERVABILITY STACK — ONE-LINE VIEW

Audit & Feedback	What did the AI say? — Raw interaction capture & user feedback
RAG Quality Metrics	Did it retrieve the right knowledge? — Retriever performance monitoring
Session Tracing	Why did the agent do that? — Full session replay with reasoning steps
Platform Tracing	What did the platform do? — Flow/Apex execution timing & errors (dev tool)
Agent Optimization	What patterns exist across thousands of sessions? — AI-driven intent analysis
Health Monitoring	Is something breaking right now? — Real-time alerting & dashboards

CHAPTER 2

Feature Comparison & Course Setup Flow

Understanding which feature does what — and why the course skips the legacy toggle

Why this matters: The Agentforce observability umbrella has four distinct features. Students often confuse them — especially because "Agent Analytics" is the name of both the old legacy feature and the new dashboard. Knowing the difference prevents setup mistakes in the org.

Feature	Purpose	Primary Audience	Status
Agent Analytics (Legacy)	Original usage & engagement data collection in Data Cloud	Admins / Analysts	⚠ Retiring May 2026
Agentforce Session Tracing	Full session capture — provisions STDm in Data Cloud. Master toggle. Everything depends on this.	All roles	✅ GA — enable first
Agent Platform Tracing	Low-level OpenTelemetry spans for Flow, Apex, LLM calls. Pinpoints performance bottlenecks.	Developers	✅ GA (Apr 2026)
Agent Optimization	AI-powered intent analysis & quality scoring. Auto-enabled with Session Tracing.	Admins / Analysts	✅ GA — auto-on

Course Setup Flow — Why It's Built This Way

- STEP 1** **Set Up Your Instance** — Enable Agentforce Session Tracing (starts capturing data). The legacy Agent Analytics toggle is intentionally skipped — this course teaches the new stack only.
- STEP 2** **Exercise 1, Task 3** — Install Tableau Next analytics app packages for Service Agent and Employee Agent. These are the pre-built dashboards that visualize the session data captured by Session Tracing.
- THE FLOW** Session Tracing captures the data → Data Cloud stores it via the STDM → Tableau Next dashboards surface it visually. Without the packages, there's nothing to view the data in.

Common student confusion: "Why don't we turn on Agent Analytics?" — Because the legacy toggle shares a name with the new experience but is a different feature entirely. This course only uses Session Tracing → STDM → Tableau Next. The legacy toggle is not part of that stack.

CHAPTER 3

Agent Analytics Dashboard — Overview Tab

High-level health check: adoption, effectiveness, and customer behavior at a glance

What this tab answers: Are users actually engaging with the agent? Is it resolving issues or are people just leaving? This is the first tab an admin or analyst should check — it tells you whether the agent is working at a basic level before you go deeper.

Global Filters

AGENT / TYPE	Filter by specific agent or use case. Use this to compare performance across multiple agents and identify which ones are underperforming.
TIMEFRAME	Adjust the reporting window. Critical for measuring whether changes you made actually improved performance — always compare before vs. after.
CHANNEL	Filter by Web, Mobile, Messaging, etc. Customer behavior differs significantly by channel — an agent that works well on web may struggle on mobile.
SUBAGENT / ACTION	Filter by specific topic or workflow. Isolates performance for one part of the agent experience without noise from others.

Performance Overview Metrics

These five metrics tell the core story of agent effectiveness.

94.12%

ENGAGEMENT RATE

Users actively interacting with the agent. High = strong adoption and willingness to engage.

0%

ESCALATION RATE

Conversations transferred to a human. 0% could mean strong automation — or a broken escalation path.

0%

DEFLECTION RATE

Issues resolved without human involvement. 0% means no successful self-service resolutions.

73.53%

ABANDONMENT RATE

Users who left before resolution. High rate signals poor responses, confusion, or unmet expectations.

—

AVG. TIME TO DEFLECTION

How fast the AI independently resolves issues. No value = no successful deflections occurred.

READING THE STORY — WHAT THESE NUMBERS TELL YOU TOGETHER

High engagement (94%) + high abandonment (73%) + zero deflection = users are willing to try the agent, but it's not resolving their needs AND it's not routing them to a human either. They're engaging and then just leaving. The agent is broken in a very specific way: it's discoverable but not functional.

High abandonment → Drill into the Sessions tab to find the intents that led to abandonment. Build new topics to address the gaps.

Low deflection → Topic instructions or actions may be too narrow — expand coverage so the agent can actually resolve requests.

0% escalation + high abandonment → Users may want to escalate but can't. Review whether the escalation path is configured correctly.

Fast deflection time → When the agent works, it works fast. Use that as a benchmark for what good looks like once you fix the gaps.

Agent Usage Overview

Bar chart showing session volume per agent — identifies most-used and highest-risk agents.

WHAT IT SHOWS Session counts broken out by agent. Helps identify which agents are getting the most traffic and which ones have disproportionate abandonment or escalation.

NOT_SET FLAG Sessions appearing as "NOT_SET" (no agent name) signal a data quality issue — likely agents built on the wrong template, or Session Tracing wasn't active before those sessions ran.

HOW TO USE THIS

Sort by highest volume. Focus your optimization effort there first — fixing a high-traffic agent with low deflection has 10x the impact of fixing a low-traffic one. Investigate NOT_SET sessions by verifying agent template and confirming the Data Kit was fully deployed before sessions began.

CHAPTER 4

Agent Analytics Dashboard — Quality Tab

AI-generated quality scoring — evaluating response accuracy, relevance, and completeness by agent, topic, and intent

What this tab answers: When the agent responds, how good is the response? The Overview tab tells you whether the agent is working. The Quality tab tells you *how well* it's working when it does. These are different questions with different fixes.

Quality Score Scale

SCALE Quality scores use a **1–5 scale** generated by AI analysis of each interaction.

Score Range	Rating	What It Suggests
0 – 2	Very Low	Responses are inaccurate, irrelevant, incomplete, or misaligned with user intent
2 – 3	Low	Partially helpful but with significant gaps — prompts, grounding, or actions need work
3 – 4	Good	Mostly relevant and accurate — targeted refinements will push it higher
4 – 5	High	Strong response quality — consistent, accurate, and well-grounded

Quality Overview — Scores by Agent

Example scores from a real org — used here as a teaching artifact.

Agent	Avg Quality Score	Rating	What It Means
Demo_HelloWorld	5.0	High	Likely a scripted test — not a reliable production benchmark
Pronto_Service_Agent	1.52	Very Low	Real production agent — responses are consistently poor quality
NOT_SET	1.0	Very Low	Sessions not mapped to an agent — data quality issue to investigate

Overall average: 1.66 / 5 — Very Low. This org has a widespread quality problem across all production agents. The "5" on HelloWorld is a scripted test, not a real signal.

Data Explorer — Subagents (Topics)

Breaks down session count and quality score by topic — the most actionable view on this tab.

Subagent / Topic	Sessions	Quality Score	Priority
Verify_Customers	45	1.47 — Very Low	Fix first — highest volume + lowest quality
NOT_SET	37	1.0 — Very Low	Investigate routing & Session Tracing setup
Order_Issues_Refunds	20	1.2 — Very Low	Fix second — high impact on customer satisfaction
escalation	6	3.67 — Medium	Agent performs best when handing off to a human
Support_Policies_and_Terms	1	1.0 — Very Low	Low volume — fix after higher-traffic topics

THE STORY THESE NUMBERS TELL

The agent scores highest when it escalates to a human (3.67) — not when it resolves on its own. Every self-service topic is Very Low. This means the agent is engaging users but can't actually help them. The escalation path is the only thing working. This is a grounding and instruction problem, not a structural one.

Very Low scores across topics → Review topic instructions and grounding data. The LLM may not have enough context to respond well — check Knowledge article coverage.

High volume + low quality → Verify_Customers (45 sessions, 1.47) is the highest-ROI fix. Start there.

Escalation scores highest → Consider routing more intents to a human until the agent is better trained. Don't hide the escalation path.

NOT_SET subagents → 37 sessions with no topic name means missed routing. Check topic configuration and confirm Session Tracing was live before those sessions ran.

Use Custom Scorers (Beta) → If the default AI scoring doesn't reflect your specific use case, define your own scoring criteria to get more relevant quality signals.

Data Explorer — Intents Tab

WHAT IT IS

Analyzes quality broken down by the specific customer intent — what the user was actually trying to accomplish.

WHY IT MATTERS

Two users can use the same topic but have completely different intents. This tab tells you which specific requests are succeeding and which are failing — a level of precision the subagent view can't give you.

HOW TO USE THIS

Filter to your lowest-quality topics, then drill into Intents to find the specific requests causing failures. For each failing intent, check: Are the session traces showing a retrieval failure? A failed action? A misclassified request? Each root cause has a different fix.

CHAPTER 5

Optimization Insights Dashboard

Strategic view — latency, quality distribution by subagent and intent, top and bottom performers

What this dashboard answers: Where exactly is the agent failing — and which failures matter most? This is the prioritization view. It connects volume, quality, and latency so you know where to invest optimization effort first.

Top Summary Metrics

26**TOTAL SESSIONS**

Overall interaction volume — context for all other metrics.

4.3s**AVG. AGENT LATENCY**

Time for agent to respond. Higher latency increases abandonment and frustration.

Very Low**AVG. QUALITY SCORE**

Overall AI quality rating. "Very Low" = responses are inaccurate, incomplete, or misaligned.

Quality Distribution by Subagent

WHAT IT SHOWS A breakdown of how quality scores are distributed across subagents — not just averages, but how many interactions fell into each quality tier (Very Low / Low / Good / High).

WHAT TO LOOK FOR If nearly all interactions for a subagent cluster in "Very Low," it's a systemic problem — not an edge case. That subagent needs a full review, not a minor tweak.

EXAMPLE FROM THIS ORG

Nearly all interactions fall into the "Very Low" category across all subagents — indicating a widespread quality issue, not isolated failures. The agent instructions, grounding, and retrieval configuration all need review together.

Top and Bottom Ranking Subagents

TOP RANKING	"No Data" — no subagent currently meets the threshold for positive ranking. This means there's no baseline to protect — every subagent needs work.
BOTTOM RANKING	Verify_Customers, Order_Issues_Refunds, NOT_SET — all Very Low quality scores. These are your immediate priorities.

Bottom rankers are the roadmap. Tackle in order of volume — highest sessions first, since that's where you'll have the most impact.

For each bottom ranker: Review topic instructions → check Knowledge article coverage → verify action configurations → test in Agent Builder before re-deploying.

No top rankers → Don't skip this. It tells you there's no "what good looks like" to model from. You're building from scratch.

Quality by Intent — Top and Bottom Ranking

WHAT IT SHOWS	Quality distribution broken down by specific user intent — the most granular view of where the agent succeeds and fails.
BOTTOM EXAMPLE	"View All Cases" — Very Low quality. The intent may be returning incomplete responses, failing actions, or misunderstanding the request.

For "View All Cases" → Pull session traces for that intent. Check: does the action return data? Is the response template complete? Is the intent being classified correctly?

Most intents Very Low → This is usually a grounding problem. The agent doesn't have enough quality context in its Knowledge base to form good answers for most requests.

Quality trending down over time → Correlate the drop with recent agent changes or new deployments. Use the time-series view to find the exact moment it degraded.

CHAPTER 6

Analytics — Service Agent Deep Dive

Six tabs — Effectiveness, Usage, Quality, Health, Trust, Voice — each measuring a different dimension

What this section covers: The full Tableau Next Service Agent dashboard. Each tab measures a distinct dimension of agent performance. Together they give a 360° view — from customer engagement down to voice interaction quality. Know which tab answers which question.

TAB 1 Effectiveness Analytics

How successfully is the agent resolving and deflecting conversations?

36.27%**DEFLECTION RATE**

Resolved without a human.
Higher = less support workload.

18.95%**ESCALATION RATE**

Handed off to a human. Higher may indicate workflow or knowledge gaps.

41.83%**ABANDONMENT RATE**

Left before resolution. High rates signal friction or poor experience.

32.35%**ENGAGEMENT RATE**

Meaningful customer interaction. Higher = stronger adoption and trust.

13.01%**SUCCESS RATE**

Workflow or conversation completion. Low rates highlight optimization opportunities.


HOW TO APPLY THIS

The Effectiveness Trend Chart shows these metrics over time — use it to see whether changes you made to the agent improved outcomes. The Session Outcome chart shows how conversations resolve (abandoned / deflected / escalated / not set) — a high "not set" count means your outcome tracking isn't fully configured.

TAB 2

Usage Analytics

How often are users interacting, and are they coming back?



HOW TO APPLY THIS

Use the Usage Trend Chart to identify adoption growth or decline. A drop in unique sessions after a deployment could mean users lost confidence in the agent. High avg. interactions/session combined with high abandonment means users are trying hard but not succeeding — a quality and content gap, not an adoption problem.

TAB 3

Quality Analytics

Are AI-generated responses accurate, relevant, and trustworthy?

AVG QUALITY SCORE	Overall AI response quality on a 1–5 scale. Lower scores indicate prompt, retrieval, or knowledge grounding gaps.
RAG METRICS	Measures how well the retrieval pipeline is grounding responses in enterprise data. Weak RAG = hallucinations or low-relevance answers even when the agent "seems" to respond.

HOW TO APPLY THIS

If Quality Score is low but Engagement is high, users are actively trying but getting bad answers. The fix is in the Knowledge base and prompt design — not the agent structure. Use the Data Explorer (Subagents and Intents tabs) to find the specific topics causing the lowest scores.

TAB 4 Health Analytics

Is the agent stable, fast, and operationally reliable?

0%**ERROR RATE**

How often interactions fail. 0% = stable workflows and reliable integrations.

279.8s**SESSION DURATION**

Average session length. Longer may indicate deep engagement or complex unresolved issues.

0.07**INTERACTION LATENCY**

AI response speed. Lower = better user experience and lower abandonment risk.

HOW TO APPLY THIS

Use the Health Trend Chart to detect reliability issues after deployments. Long session duration paired with high abandonment = users are spending a lot of time with an agent that never helps them. That's a quality problem, not a performance problem.

TAB 5 Trust Analytics

Is the agent following its instructions and producing safe, compliant responses?

INSTRUCTION ADHERENCE

How consistently the AI follows configured instructions and guardrails. Higher = more reliable and controlled behavior. A drop here after a prompt change is a red flag.

TOXICITY SCORE

Measures potentially harmful or inappropriate responses. Lower scores improve trust, compliance, and user safety. Even a small uptick warrants immediate investigation.

HOW TO APPLY THIS

Trust metrics are your governance layer. Monitor Instruction Adherence after any prompt update or Knowledge base change — if it drops, the agent may be ignoring guardrails. Use Toxicity Score to identify edge cases where the agent generates responses that violate your content policies before they reach large user volumes.

TAB 6

Voice Analytics

How smooth is the conversational experience in voice-based AI interactions?

0%**INTERRUPTION RATE**

How often users interrupt the AI mid-response. 0% = smooth, natural interactions. Higher rates may indicate delays or inaccurate responses prompting user frustration.

HOW TO APPLY THIS

Voice is a different UX contract — users expect near-instant responses and natural pacing. A rising interruption rate is often a latency problem, not a content problem. Cross-reference with the Health tab's Interaction Latency. If latency increases and interruption rate follows, the fix is performance optimization, not prompt rewriting.

CHAPTER 7

Optimization — Tagging & Intent Tags

How Agentforce surfaces patterns across thousands of sessions without you reading each one

What this covers: Tagging is the mechanism Agent Optimization uses to make sense of session volume at scale. Without tags, you have a list of transcripts. With tags, you have a map of what your users actually need — and where the agent is failing to deliver it.

How Tagging Works — The Flow

- STEP 1** Each session is broken into **moments** — one per distinct user intent within the conversation.
 - STEP 2** Each moment receives an **Intent Tag** — an AI-generated label describing what the user was trying to accomplish (e.g., "Reset Password", "Check Order Status", "Verify Account").
 - STEP 3** Tags with enough volume are grouped into **System Tag Clusters** — cross-session patterns that reveal what users consistently ask about, regardless of how they phrase it.
-

Critical Rules — What Students Must Know

Rule	Detail	Why It Matters
Clusters are generated weekly	Not real-time — batch process runs once per week	Don't expect instant results after enabling Session Tracing
Minimum 50 moments required	Clusters won't appear until 50+ moments exist	Low-traffic orgs may not see tags for weeks
Minimum 30 days of data	30-day window required before clusters generate	New orgs need patience — the system is learning
10+ moments per cluster	A cluster only gets a tag if 10+ moments are associated	Very rare intents won't surface as named clusters

Most common student/customer question: "Why don't I see any tags yet?" — Answer: You need enough volume (50+ moments) and enough time (weekly batch, 30-day minimum). Tags won't appear immediately after enabling Session Tracing. Set expectations with customers during implementation.

How to Use Tags Once They Appear

Find the highest-volume clusters → These are what your users care about most. If they don't have a mapped topic in the agent, build one.

Find clusters with low quality scores → These are intents the agent is attempting but failing. The tag name tells you the intent; the quality score tells you how badly it's failing.

Find clusters with no topic match → These are gaps — real user needs the agent can't currently address. Prioritize building topics for the highest-volume unmatched clusters.

Track clusters over time → A cluster that grows week-over-week and stays Very Low quality is a compounding problem. Fix it before it becomes your top abandonment driver.

FULL OBSERVABILITY STACK — END-TO-END REFERENCE

Setup Toggles	Turn on the right layers — Session Tracing is the master switch everything else depends on
Overview Tab	Is the agent working? — Engagement, abandonment, deflection, escalation at a glance
Quality Tab	How good are the responses? — AI-scored quality by agent, topic, and intent
Optimization Insights	Where exactly is it failing? — Latency + quality + top/bottom performers for prioritization
Effectiveness Tab	Is it deflecting, escalating, or abandoning? — Automation effectiveness over time
Usage Tab	Who is using it and how much? — Adoption, session depth, engagement trends
Health Tab	Is it operationally stable? — Error rate, latency, session duration
Trust Tab	Is it safe and compliant? — Instruction adherence and toxicity monitoring
Voice Tab	How smooth is the voice UX? — Interruption rate and conversational flow
Intent Tags	What do users actually need? — AI-clustered patterns that reveal gaps and priorities